

Goal:

To understand, use and analyze performance of synchronization mechanisms in Linux kernel. In this assignment, we will use various lock implementations offered by kernel and add custom implementation for a read-write lock. Two parts of the assignment are explained below.

PART 1

Get a handle of the sample code (attached), go through the README and understand the code organization. Follow the steps to compile the provided module and execute user-space micro benchmark (syncbench) to understand its usage and output formats. The sample code and template implements the in-built kernel spinlock synchronization mechanism. As part of the assignment you are required to extend the kernel module to implement the following,

- (i) use kernel API to use read-write locks, sequential locks and RCU
- (ii) design a custom read-write lock and use it just like using a kernel provided locking scheme. Verify both the above implementations for correctness.

PART 2

Perform an empirical comparison of all locking schemes (implemented in PART 1). The comparison should focus on overheads considering read-write ratio as a parameter for comparison. Following guiding principles may be helpful in empirical analysis,

- (i) Clear explanation of experimental setup, workload and parameters etc.
- (ii) Explanation of results, verifying statistical sanctity (run-time, avg., deviation etc.)
- (iii) Result insights

Extra credit

Design a new locking scheme to improve performance beyond the best locking scheme found in the empirical analysis (PART 2).

Submission format

Submissions should be one single archive (named asg2-rollno.tgz) with the following structure

when expanded.

asg2-rollno

```
|
|-----src      source code, Makefile and an implementation notes text file
|-----run
|               |----- correctness.txt      Output of tests showing correctness of your
|               |                           implementation
|               |----- unit-test.sh (or .py etc.)  when executed, it will execute all the
|               |                           correctness tests reported in correctness.txt
|               |-----analysis-test.sh (or .py etc.) when executed, it will execute all tests which
|               |                           generates data used to write your empirical
|               |                           analysis report inside doc folder
|-----doc
|               |----- report.pdf
|               |----- plots and figures
```

The report should be a document which enables the reader---to understand your experiments and the findings, to re-run the experiments if she desires so.

"Do not defeat your purpose of being here by cheating"

IITK anti-cheating policy: <https://www.cse.iitk.ac.in/pages/AntiCheatingPolicy.html>